





Database World



	→ Big 3 Cloud	Prop. → Proprietary
SaaS	→ SaaS	
OSS	→ Open Source	

Traditional DB




EXAMPLES

OSS Postgres
 Prop Oracle
 OSS MySQL
 AWS RDS
 SaaS Yugabyte
 SaaS TiDB

CONCEPTS

- WAL & Journaling
- Buffer Pool
- Mem. \rightleftharpoons Disk
- ACID, BASE
- CAP & PACELC
- Distributed databases
 - Txn & Isolation levels
 - MVCC
 - Query execution
 - Master-Slave
 - Master-Master (Active-Active)
 - 2PC vs Quorum \rightarrow Paxos, RAFT
 - Storage engines
 - \hookrightarrow MySQL (InnoDB, MyISAM...)
 - \hookrightarrow RocksDB
- Clock Synchronization
- SQL flavour
- Materialized Views, Views
- Window fn.
- Connection Pool
- Btree vs Log Structured Merge Tree



Columnar Storage







(for OLAP & Data Warehouse workloads)

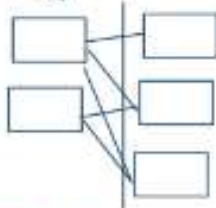




EXAMPLES





- oss Apache Cassandra (datastax)
- oss Apache HBase
- ☁ Google BigTable
- ☁ Snowflake
- oss Clickhouse
- oss Apache Hive
- oss Apache Impala
- oss Presto (Query engine)
- ☁ Amazon Redshift
- ☁ Google Bigquery
- ☁ Azure Synapse
- oss Apache Pinot
- oss Apache Druid


CONCEPTS


- Tablets
- Row vs Column Store
- Vectorized execution
- Costs
- OLAP vs OLTP
- Vectorize Query execution
- SIMD H/W optimization
- complex joins
 - ↳ || SORT-Merge
 - ↳ || hash
 - ↳ multiway Worstcase
- User defined fns.

<div>  <div>Object Store</div> </div>	
EXAMPLES	CONCEPTS
<div>  AWS S3  Google Cloud Storage  Azure Blob Storage  Pure Storage  Ceph (Openstack) </div>	<ul style="list-style-type: none"> • S3 API • Metadata service • Consistency & Redundancy

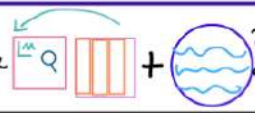
<div> <div> <div>Key</div> <div>Value</div> </div>  <div>Key-Value DB</div> </div>	
EXAMPLES	CONCEPTS
<div>  Redis  Memcached  AWS ...  TiKV </div>	<ul style="list-style-type: none"> • Caching Strategies <small>(redis, memcached, etc.)</small> • Disk persistence • SQL vs NoSQL Query <small>(memcached) (redis)</small> • Redis Internals • Copy on write • Cache vs In-memory datagrid

<div> <div>   </div> <div>Document DB</div> </div>	
EXAMPLES	CONCEPTS
<div>  MongoDB  AWS DocumentDB </div>	<ul style="list-style-type: none"> ◦ ACID, Txn, JOINS later support ◦ mmap → WiredTiger (mongo)

 Graph Database	
EXAMPLES	CONCEPTS
OSS Neo4j Amazon Neptune	<ul style="list-style-type: none"> Native Graph Processing & Storage Transaction Handling

 Vector DB	
EXAMPLES	CONCEPTS
SaaS Pinecone SaaS Weaviate OSS Chroma OSS Traditional "DB with vector" ↳ Postgres ↳ Cassandra ↳ Clickhouse	<ul style="list-style-type: none"> Vectors Embeddings Distance metrics KNN vs ANN Index <ul style="list-style-type: none"> Flat Index Inverted File Index

 Time-Series DB	
EXAMPLES	CONCEPTS
OSS Graphite OSS OpentsDB SaaS InfluxDB (Pulsar) OSS Prometheus (Pulsar)	<ul style="list-style-type: none"> Sampling Rate Querying

 Data Lakehouse	
EXAMPLES	CONCEPTS
OSS dremio OSS Deltalake	<ul style="list-style-type: none"> File formats: Parquet, Arrow etc. Integration with OLAP, BI etc. Data governance